Break Before You Shape

Exact line breaking from a precompiled, certified width model

Draft · 2026-11-03 · every number in this paper is read from result files produced on one Apple M3 machine and on GitHub Actions

Abstract. Text engines shape a paragraph before they can break it into lines, because widths depend on kerning, ligatures and contextual forms. We show that, for the fonts and scripts we test, a line's shaped width is a *local* function of its text, and that this locality can be *proven* from a font's layout tables. We compile a shaper's behaviour into a small per-font model by measuring it as a black box, and a static analysis of OpenType and AAT tables certifies the lines on which the model is exact. Lines are then broken before anything is shaped. Across 449 local font faces, 5,339,408 certified lines (96% of all) contain no error. Inside Flutter, a 10 KB (gzipped) file built once on the developer's machine reproduces Flutter's own line breaks and paragraph heights in 2,524,902 of 2,524,902 paragraph layouts on macOS (nine corpora in six languages, six font sizes, twelve widths), and in all 5,529,364 it decides on Linux, macOS, an iOS simulator and an Android emulator in continuous integration. The same holds with letter and word spacing, height multipliers, `maxLines` and hard line breaks (4,575,628 layouts), and for Hebrew, mixed-direction, Arabic, Japanese, Chinese and Korean text in four more fonts (1,209,076 layouts): not one layout differs. The breaker hands 0.87% of Latin-script layouts back to the engine and computes paragraph heights 18–25× faster than `TextPainter`.

1 · Introduction

Every mainstream text stack (Chrome, WebKit, Android, Skia's SkParagraph, CoreText, DirectWrite) runs the same order: *shape* the paragraph with a shaper such as HarfBuzz, *break* it into lines using the shaped advances, and often *reshape* around the breaks. Shaping is the expensive step, and much of it is wasted. A chat list computing thousands of item heights, or an editor re-wrapping on every resize, shapes text that it never draws. Figure 1 shows the alternative this paper makes exact: break first, from a precompiled model, then shape only what is drawn.

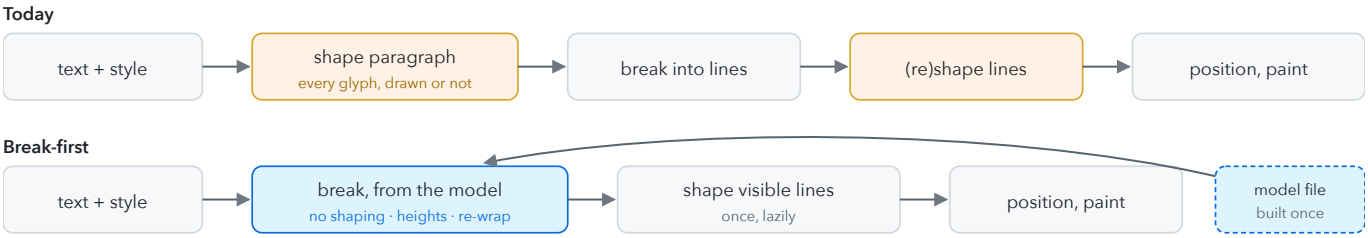


Figure 1. Shape-first layout shapes every glyph before it knows where lines end. Break-first computes line breaks and heights from a precompiled model, so text that is measured but never drawn is never shaped.

Contributions.

1. **Distillation.** A shaper is measured as a black box (every character, every pair, every ligature island in context) and compiled into exact tables. The model matches the target engine by construction.
2. **A certificate.** A static analysis of OpenType (GSUB, GPOS, GDEF) and AAT (`morx`, `kern`, `kerx`, state machines included) tables proves on which lines the model's width equals the shaped width, including Arabic letters per joining form.
3. **Break-first layout.** Width-independent records give every candidate line's width in O(1), so a resize costs one greedy pass. The breaker reproduces the host's measured conventions (spacing, line heights, hard breaks, bidirectional runs, split kerning) and defers wherever the host's behaviour is uncertain.
4. **Evidence.** A census of 449 font faces, and a Flutter package checked against Flutter's own layout on four platforms, for Latin, Greek, Cyrillic, Hebrew, Arabic, Chinese, Japanese and Korean text, with five fonts.

2 · Width is local, and can be distilled

A line's shaped width is the sum of its glyph advances plus positioning adjustments. In practice those adjustments are *local*: pair kerning touches two adjacent glyphs, and a ligature replaces a short run of letters. So the width of any line decomposes into a few kinds of terms: one advance per character, one term per adjacent pair, and a term for each ligature *island* with its two neighbours. Figure 2 shows a real decomposition, measured from Flutter.

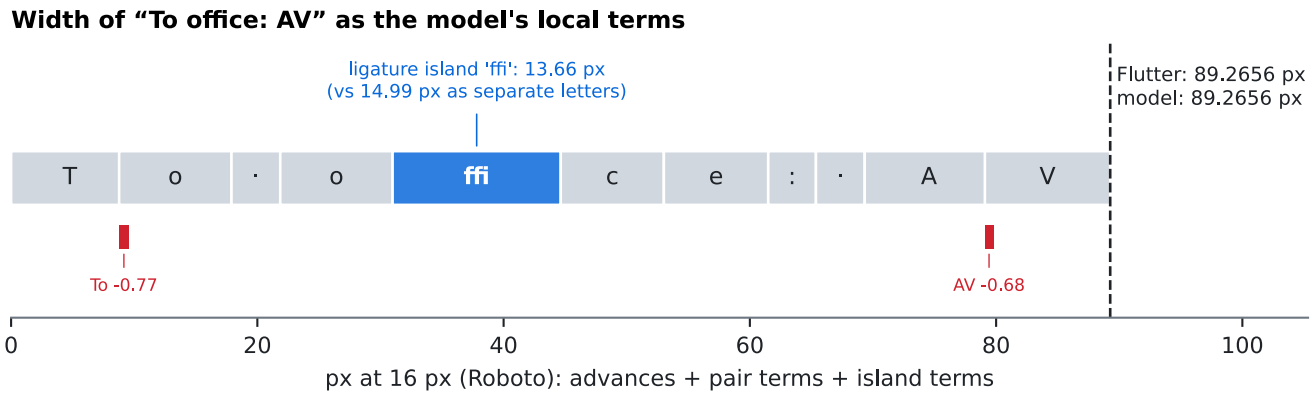


Figure 2. “To office: AV” in Roboto at 16 px. Gray: advances alone, which sum to 92.055 px. Red: pair terms (kerning). Blue: a ligature island. The model sums to 89.2656 px; Flutter measures 89.2656 px.

Distillation. We do not parse kerning or ligature tables. We ask the target shaper. For Flutter, the oracle is Flutter's own `ui.Paragraph`, run under `flutter_tester` on the developer's machine:

- every character alone;
- every ordered pair;
- every certificate island alone, and next to every safe neighbour (§3).

We measure at a font size equal to the font's units per em, so values are font units. All 614,675 measurements came back as exact integers, and widths scale linearly with size (§5.5), so **one file serves every size**. For Roboto (Latin, Greek and Cyrillic, 699 characters), identical rows and columns of the pair matrix collapse into 114×121 classes. With 90 islands, the file is 577 KB, 10.3 KB gzipped, and builds in 8.8 s.

---

3 · A certificate from the font's tables

Locality is not guaranteed: a contextual substitution can make a glyph depend on text arbitrarily far away. So each character is classified from the font's layout tables, without running the shaper:

- **S** (safe): no substitution can change or consume its glyph;
- **P** (primer): it can start a substitution, such as the first letter of a ligature;
- **K** (passive): it can only be consumed by a substitution a primer started;
- **U** (unsafe): everything else.

An *island* is a maximal run of P/K characters that contains a primer.

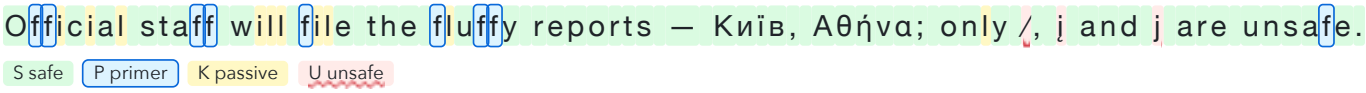


Figure 3. Roboto's certificate on sample text. Islands form only around *f*. Unsafe characters send their paragraph to the engine.

> **Theorem.** Let a line contain no U character, and let every island in it have at most m characters. Then the shaped width of the line equals the model's width.

Proof sketch.

1. **Substitution leaves S glyphs untouched, and islands are closed.**
   - *OpenType.* A single, multiple or alternate lookup fires only at its coverage glyphs, which are U. A contextual rule fires only if every position matches a glyph that can occur, so for each live rule the cheapest position's characters are made U. Ligatures start at P and consume only P/K glyphs.
   - *AAT.* Every `morx` subtable (rearrangement, contextual, ligature, insertion, in either processing order) and the `kerx`/`kern` state machines are emulated as HarfBuzz runs them. An S class changes nothing from any reachable state and returns to a state equivalent to a start state (state minimisation). Every island up to the bound runs through the whole chain from each start state, with every mark and stack entry from before it stale; if it touches stale state or depends on its surroundings, its cheapest character becomes U.
2. **Positioning is pair-additive.** Only single adjustments, PairPos with ValueFormat2 = 0, the `kern` and `kerx` pair formats, and kerning machines proven local by step 1 can fire.
3. **Width splits into local terms.** By 1 and 2, the width is a sum of terms, each touching at most one island and its S neighbours.

4. **The terms agree.** The model's terms are measured values of exactly those terms. Islands are measured with every S neighbour, and the checker verifies no model term spans two islands. ∎

The proof trusts that the shaper implements OpenType and AAT as specified, that fontTools parses the tables correctly, and the analyzer's code, which is tested but not machine-checked; its AAT emulation was compared with `hb-shape` on about 22,700 islands. By default it also trusts that the shaper selects lookups as HarfBuzz does (the scripts certified, language, default variation coordinates); a strict mode drops that assumption.

Two additions for the breaker. *Pair roles:* for each character, whether its glyph can be the first or second glyph of any pair adjustment (GPOS pair and cursive coverage, `kern` pairs; anything not analysed counts as every glyph). Only first × second pairs are measured, so a CJK font with 28,092 characters needs no pair table at all and one with 16,644 builds in under a second. *Joining forms:* each Arabic (letter, form) is certified as a virtual character whose glyph is the font's isol, init, medi or fina substitute; the breaker picks forms from Unicode joining types, as HarfBuzz's Arabic shaper does, so lam-alef is just an island.

4 · Breaking before shaping, the way the host does

For each paragraph, one pass builds a width prefix in units, and a second finds break opportunities. Both are independent of the box width, so re-wrapping at a new width only re-runs the greedy pass over these records. To reproduce a specific engine, the breaker must also reproduce that engine's conventions. We measured Flutter's rather than assuming them (Table 1), and the breaker **defers** whenever a decision depends on something it cannot certify.

Flutter (SkParagraph) convention	How it was established	Breaker
Shapes once, breaks on shaped cluster advances; kerning sits on the left glyph	Exactness against Flutter's breaks	Pair terms attach to the left character
Line fits if its width, floored to 1/100 px in float32, ≤ box	Accepted and rejected widths around the box edge	Same test, emulated in float32
Line height = font line height, rounded to whole px	Every height in Table 2 (2,524,902 on macOS)	lines × rounded height
No kerning across script runs; spaces join the preceding run	"Time Machine" inside Russian text	Pair term 0 at a script change
UAX #14 with Unicode 15.1 number rules; word-initial dash binds	Per-class tallies of taken vs refused breaks	Rule table for the characters it handles
Sometimes breaks at a space inside an unbreakable segment	Examples (e.g. `self ⏎!= 1.`)	Defers when such a space could end the line
Positions drift in float32 at some sizes	Drift by size (§5.5)	Defers decisions the drift could flip
Android truncates the font size to 1/64 px (e.g. 18.4 → 18.390625); macOS and iOS do not	Effective size implied by measured widths at 8 sizes, on the emulator	One probe layout at start-up detects it; sizes are truncated the same way
`\n`, U+2028 and U+2029 end a line; the break kerns with the character before it as a space does on macOS and iOS, not on Linux or Android	"LT↵" loses Roboto's T-space kern on macOS only; CI on Linux	Mandatory records; a start-up probe decides the kerning
Letter spacing is added to every glyph advance (half before the first); positive spacing turns ligatures off, negative keeps them	Widths with spacing against widths with `liga` off	A second certificate without ligatures; spacing per glyph cluster
Word spacing goes on the last space of each run that is a break opportunity, also after a hard break	Glyph positions with and without spacing	Counted per record; defers where a space meets a bidi run boundary
With a height multiplier, a line is round(size × height), size and height taken as float32	15 × 1.3 gives 19, not 20 (heights in Table 4)	The same formula
Kerning from `kern`/`kerx` tables is split between the two glyphs	Times New Roman's space-A kern at a line start	Each pair's right-glyph share, measured from glyph boxes
Bidirectional text: UAX #9 levels; shaping runs split at level changes, and scripts are itemized within each run	Hebrew and Arabic UI strings	Each run measured with the sidecar distilled in its direction
UAX #14 rules LB15c, LB20a, LB21a/b, and LB19a's and LB30's East Asian exceptions; small kana break like ideographs	Validation mismatches, one rule each	Rule table extended

Table 1. Host conventions the breaker reproduces, each measured from Flutter.

5 · Evaluation

5.1 · Census: how often does the certificate apply?

We ran the analyzer and the checker on every font face on the machine that covers basic Latin (449 faces). For each, the checker compares every certified line of a 1,000-paragraph English corpus sample at three widths against HarfBuzz, using exact ground truth that reshapes every candidate line.

- **Errors:** 5,339,408 of 5,560,372 lines (96.0%) were certified, and **0 were wrong**.

- **Coverage:** 410 faces certify at least 99% of their lines; 0 certify none.
- **Causes of failure** (39 faces under 99%, Figure 4b): in 24 the font's substitutions genuinely touch common letters; 6 have AAT contextual or ligature machines over common letters that are not local (Apple Chancery, Marion, Aqua Kana); in 5 the exhaustive island check's budget lowered the island bound below 4; the rest lack characters or class letters as marks.

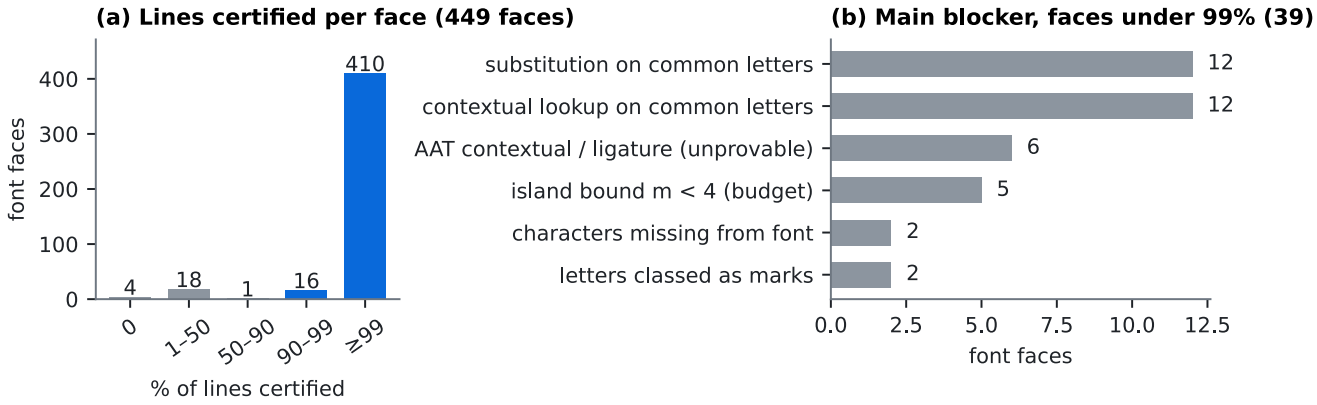


Figure 4. (a) Share of lines certified, per face. (b) The reason behind the most frequent unsafe character, for the faces under 99%.

5.2 · Today's shortcuts are not exact

The usual alternatives to full reshaping lose context, and the error is measurable. Against exact ground truth in HarfBuzz (14 fonts, 5,170 paragraphs, 6 widths; Zapfino is left out, because its long contextual swashes are exactly what the certificate hands back to the shaper):

- **Shape once, then break** (what SkParagraph does) gets up to 1.2% of paragraphs wrong.
- **Per-word caches** (Minikin, Pretext) get up to 9.6% wrong.
- **The distilled model** gets 0.0% wrong.

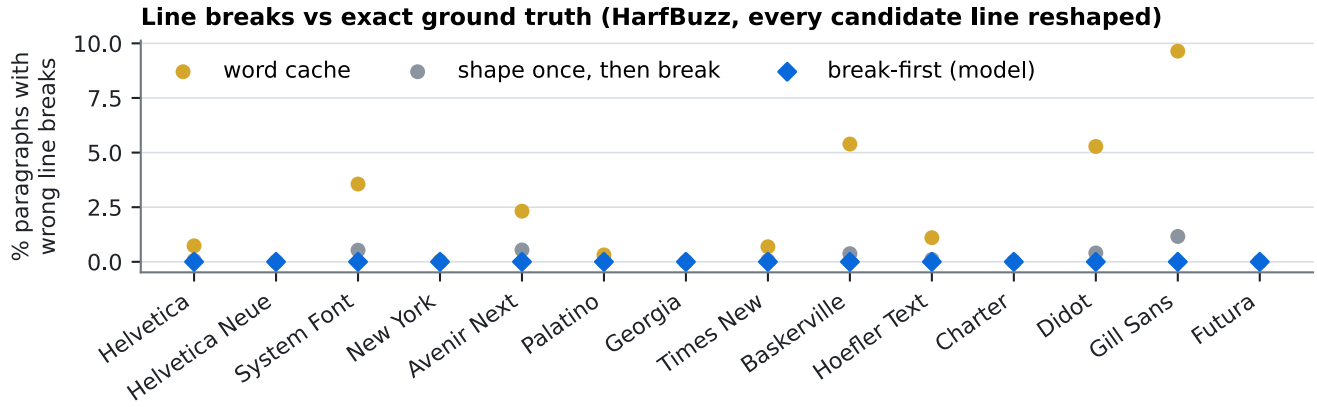


Figure 5. Share of paragraphs whose line breaks differ from exact ground truth (C prototype, HarfBuzz 15).

5.3 · Exactness inside Flutter

The real test is a production engine. A Flutter app (Flutter 3.47.5) loads the sidecar built on the host and compares every line break and every paragraph height with Flutter's own layout. The run covers nine corpora:

- English Python documentation, man pages and license texts;
- user-interface strings in English, German, French, Russian, Ukrainian and Greek, from macOS localizations.

Each paragraph is laid out at 6 font sizes (12, 13, 16, 18.4, 21.3, 24 px) × 12 box widths.

Platform	Paragraph layouts	Identical breaks	Identical heights	Deferred	Canary
macOS, release (all paragraphs)	2,547,000	2,524,902 / 2,524,902	2,524,902 / 2,524,902	0.87%	0.0 px
iOS 17.5 simulator, debug (300 per corpus)	194,184	192,193 / 192,193	192,193 / 192,193	1.03%	0.0 px
Android 15 emulator, release (300 per corpus)	194,184	192,176 / 192,176	192,176 / 192,176	1.03%	0.0 px

Table 2. One sidecar file, built once on macOS, checked against Flutter's layout on three platforms. Lines checked on macOS: 7,201,290.

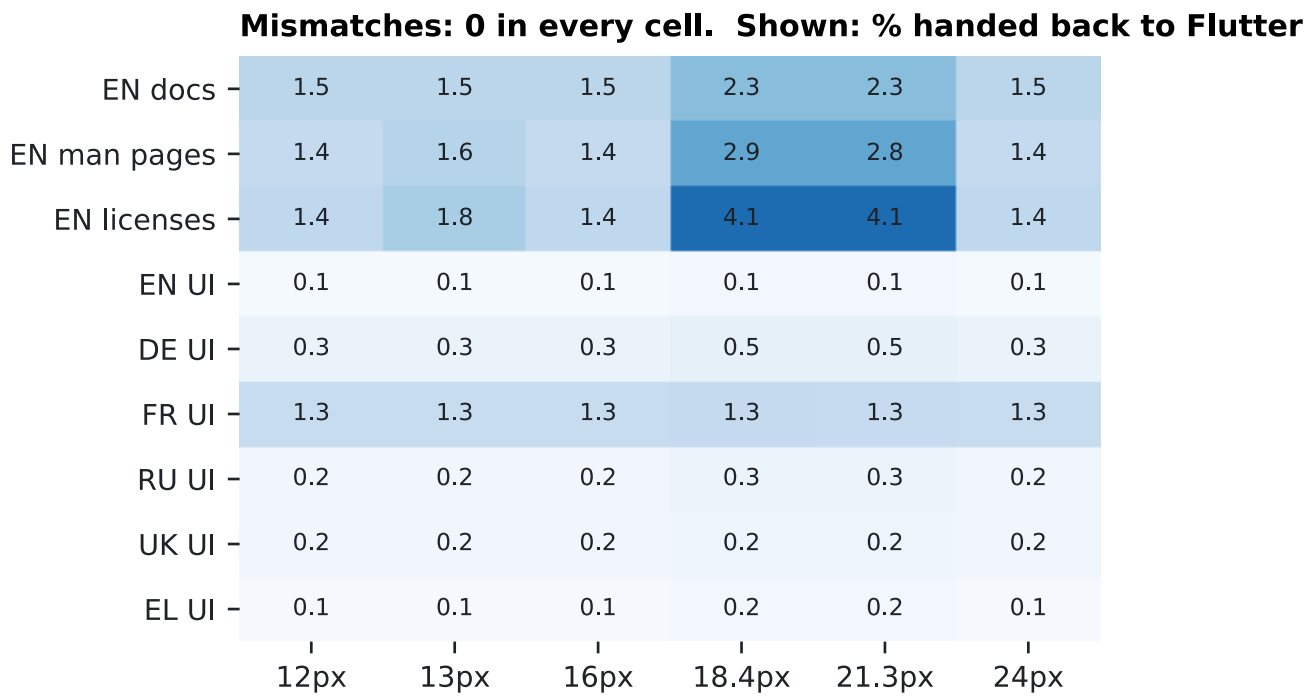


Figure 6. No cell has a single mismatch. Values are the share of paragraph layouts the breaker hands back to Flutter: words wider than the line, characters outside the model, and decisions it will not guess.

On other machines. The same checks run on GitHub Actions (run 41614792200), where each job builds its corpora from its own runner's files. The Linux engine, like Android's, rounds font sizes to 1/64 px; the start-up probe finds it.

Platform (GitHub Actions)	Identical breaks	Different	Heights equal	Size rounding
macOS 15, release (virtual Apple M1)	2,518,521	**0**	2,518,521 / 2,518,521	none
Ubuntu 24.04, Linux desktop release	2,475,396	**0**	2,475,396 / 2,475,396	1/64 px
iOS simulator (CI)	192,168	**0**	192,168 / 192,168	none
Android 15 emulator, x86_64 (CI)	190,348	**0**	190,348 / 190,348	1/64 px
flutter_tester, macOS	76,836	**0**	76,836 / 76,836	none
flutter_tester, Linux	76,095	**0**	76,095 / 76,095	1/64 px

Table 3. The shipped sidecar against Flutter on GitHub's machines, default style.

Style options and hard breaks. Each rule in Table 1 was checked on the same corpora, joined one to four paragraphs at a time with `\n` (14,115 texts), at 6 sizes and 8 widths, against Flutter laid out with the same options:

Option	Layouts	Identical breaks	Different	Heights equal	Deferred
hard breaks	677,520	657,974	**0**	657,974 / 657,974	2.9%
letterSpacing 0.5	677,520	655,457	**0**	655,457 / 655,457	3.3%
letterSpacing -0.3	677,520	646,691	**0**	646,691 / 646,691	4.6%
wordSpacing 2	677,520	657,789	**0**	657,789 / 657,789	2.9%
height 1.3	677,520	657,974	**0**	657,974 / 657,974	2.9%
maxLines 3	677,520	657,974	**0**	657,974 / 657,974	2.9%
combined	677,520	641,769	**0**	641,769 / 641,769	5.3%

Table 4. macOS release build. “Combined” sets letter spacing 0.7, word spacing 1.5, height 1.2 and `maxLines` 4 together.

5.4 · Other scripts and directions

Right-to-left and East Asian text use fonts the platform already has, so nothing is redistributed: a sidecar is built through Flutter for each and checked on 4,000 UI strings per language, from macOS localizations, at 3 sizes and 6 widths, alone, joined with hard breaks, and with spacing.

Text · font	Layouts	Identical	Different	Decided
Hebrew, right-to-left paragraphs · Arial	167,994	83,834	**0**	50%
Hebrew + Latin and digits, mixed direction · Arial	239,994	210,018	**0**	88%

Hebrew, right-to-left · Times New Roman	167,994	83,922	**0**	50%
Hebrew + Latin and digits, mixed · Times New Roman	239,994	210,397	**0**	88%
Arabic (joining forms), right-to-left · Arial	335,988	28,214	**0**	8%
Arabic (joining forms) · Times New Roman	335,988	33,526	**0**	10%
Japanese and Chinese · Hiragino Sans GB	335,988	278,143	**0**	83%
Korean and Japanese · Apple SD Gothic Neo	335,988	281,022	**0**	84%

Table 5. "Decided": the share of layouts the breaker computes itself; the rest (vowel marks, characters a font lacks, words wider than a narrow line, Arabic mixed with other directions) go to Flutter. Hiragino Sans GB has no pair positioning (28,092 characters, no pair table); Apple SD Gothic Neo's 16,644 characters need 57×51 pair classes, a 316 KB sidecar.

5.5 · Speed

All methods run on the same paragraphs, the ones break-first decides itself, at 16 px. Each time is the median of 5 passes inside one run of a macOS release build, and then the median of 3 such runs; the speedups below differ by at most 12% between runs.

- **Paragraph heights:** break-first is **18.5×** (docs) and **25.2×** (chat) faster than `TextPainter`.
- **Re-wrapping at a new width** is **11.6×** (docs) and **9.6×** (chat) faster than Flutter's re-layout, which already skips reshaping.
- **A warm word cache** is about as fast, but only after measuring every word once (the cold column), and it still breaks 3 docs and 3 chat paragraphs differently from Flutter.
- **End to end**, including paragraphs handed back to Flutter, the gain is 16.9× (docs) and 23.1× (chat).
- **Android emulator**, release, one run: 33.3× (docs) and 34.8× (chat).
- **On GitHub's machines** (median of 3 runs; one on Android): 26.0× / 27.6× on a virtual M1, 18.9× / 24.7× on Linux, and 24.5× / 38.6× on the Android emulator (docs / chat).
- **Native code** does not change the picture. A line-for-line C port called through `dart:ffi` gives the same line starts in all 2,552,760 checks and is 2.0× (docs) and 2.3× (chat) faster than the Dart breaker, 2.2× end to end; since the breaker is already a small share of the cost it replaces, the package stays in Dart, which also runs on the web.

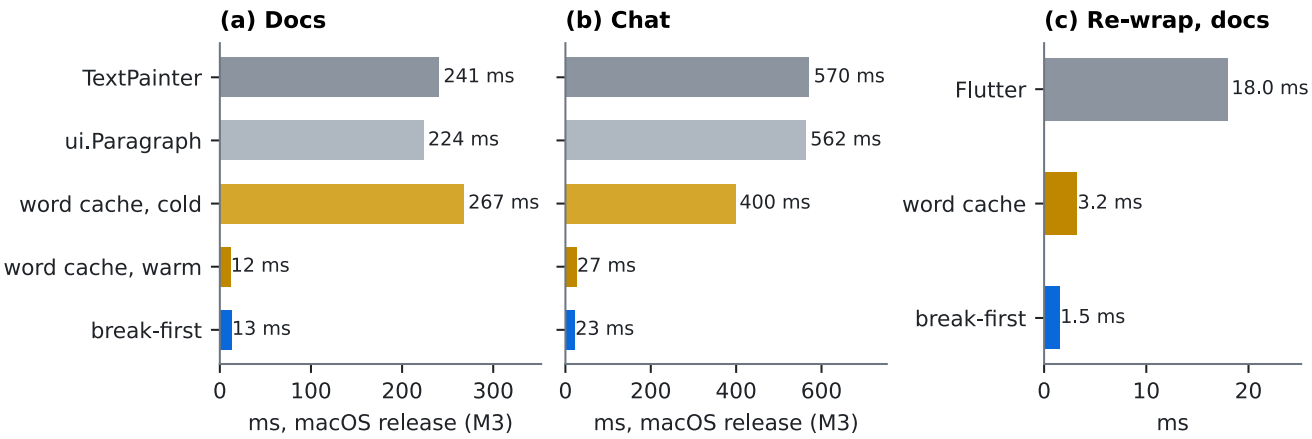


Figure 7. Time to compute every paragraph's height. (a) 5,089 Python documentation paragraphs at 400 px. (b) 23,736 UI messages at 280 px. (c) Re-wrapping the already laid-out docs paragraphs at a new width.

An app-level case: exact list extents. A chat list needs every item's height to place its scrollbar and jump to a message. For 24,000 messages, break-first computes all heights in 33 ms against 588 ms with `TextPainter` (18.0×), fallbacks included. 24,000 of 24,000 heights are equal, and the largest scroll-offset error over the whole list is 0.0 px. On the Android emulator the gain is 19.7×, again with 24,000 equal heights.

The package's example app shows what this buys (video in the repository). A stock `ListView` asked to jump to message #2,500 of 10,000 can only estimate where it is: it lands on #1,969, in a 717 ms frame spent laying out every message on the way. With break-first extents the list lands on #2,500 in a 22 ms frame, and knows its exact total height. All 10,000 extents equal `TextPainter`'s, which takes 444 ms to compute them (Android emulator, release).

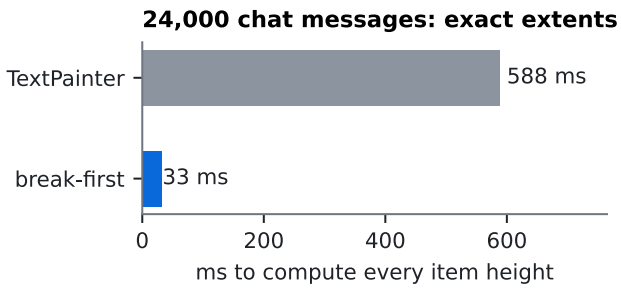


Figure 8. Computing every item's exact height for a chat list.

5.6 · Font sizes and float drift

Flutter's widths are linear in font size, so one file in font units serves every size. But SkParagraph accumulates positions across a paragraph in float32.

- **Exact sizes** (10, 12, 16, 20, 24, 32 px for Roboto): size/units-per-em is a short binary fraction and positions are exact.
- **Other sizes:** positions drift by up to 1.3e-05 of the x position, as much as 0.40 px on a paragraph-long line.

The breaker certifies exact sizes outright. At other sizes it defers any decision that a width error of 2 × 2.5e-05 of the x position could flip: drift at both ends of the line, each allowed 1.9× the largest drift measured. On Android the engine also truncates the font size to 1/64 px, which a single start-up probe detects. With both measures, all 6 sizes in Table 2 had zero mismatches on every platform.

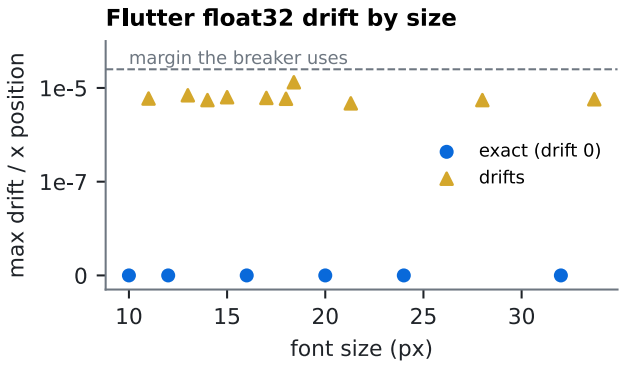


Figure 9. Largest drift between the model and Flutter, relative to x position, by font size.

6 · Related work

- **TeX's TFM lig/kern programs** are a hand-made per-font width automaton: the closest ancestor of the model, without contextual forms or a certificate.
- **PDF libraries** sum AFM widths and kerning pairs, ignoring substitution.
- **WebKit's simple text path** skips the shaper only when kerning and ligatures are off.
- **Android's Minikin and Pretext** measure and cache words, which loses context across words (§5.2).
- **HarfBuzz's** `UNSAFE_TO_CONCAT` makes reshaping after breaking cheaper, but keeps the shape-first order.
- `typeset` (2026) computes widths from advances and pair kerning with a HarfBuzz fallback, without contextual terms or a proof.

We found no prior static analysis of shaping tables. This survey is a web search, not an exhaustive review.

7 · Limitations

Everything outside this list is decided exactly or handed to Flutter; what is listed is handed to Flutter or not claimed.

- **Scripts.** Indic syllables (reordering), Thai, Lao, Khmer and Burmese (dictionary line breaking), vowel and other combining marks, emoji, and Arabic mixed with other directions go to Flutter, as does any character a font lacks (fallback fonts differ between devices).
- **Styles.** A paragraph has one font and size: rich text with several, font features and `strutStyle` go to Flutter.
- **Proof.** The trust base in §3; the AAT emulation follows HarfBuzz 15 and is tested, not proven. Fonts with only a `mort` table, rare `kerx`/`kern` formats, and islands beyond the stated bounds are blocked rather than analysed.
- **Measurement.** Mobile timings come from simulators and emulators, not phones.
- **Scope of the gain.** The speedup is for measurement and line breaking; painting is unchanged, and lines that are drawn are still shaped.

8 · Reproducibility

All code and results are in the accompanying folder:

- `certify.py`: the analyzer;
- `bfl.c`: the C prototype and checker;
- `census.py`: the census;
- `packages/break_first/`: the Flutter package (breaker, sidecar build step, tests, list example);
- `flutter_bench/`: the benchmark app, the validation and timing harness, and the script tests;
- `native/`: the C port; `.github/workflows/`: the CI that reruns the benchmarks;
- `paper/make_figures.py`: the figures and every number in this paper, read from the result files.

Environment: Flutter 3.47.5 (Dart 3.13.4), HarfBuzz 15.0.0, macOS on an Apple M3, iOS 17.5 simulator, an Android API 35 emulator, and GitHub's macOS 15 and Ubuntu 24.04 runners.